

Workshop

Introducing SQL: A Foundation of Data Analytics

Robb Sombach
University of Alberta
Alberta School of Business



Agenda

- Introduction
 - Why SQL?
 - What about Python? R?
 - Data Analytics
- Relational Database
 - What is a database?
 - Terminology
 - SQLite
 - **Exercise 1**
- SQL
 - Data Definition Language (DDL)
 - **Exercise 2**
 - Data Manipulation Language (DML)
 - **Exercise 3**
- Open Data Portal
 - How I prepared for today

Robb Sombach

- Work Experience
 - 15+ years working in the IT industry
 - 10+ years Self-Employed IT Consultant
- IT Positions
 - Systems Analyst / Business Analyst
 - Database Administrator (Oracle / SQL Server)
 - Network Administrator
 - Developer

Robb Sombach

- Teaching Experience
 - 5 years teaching at NAIT
 - Computer Systems Technology (CST)
 - Digital Media and Information Technology (DMIT)
 - 6+ years teaching at University of Alberta
 - Technology Training Centre
 - Alberta School of Business

Resources

All Workshop files can be downloaded here

http://bit.ly/odd_2019

Introduction

Workshop

Introducing SQL: Foundation of Data Analytics

Goals

- Introduce relational database concepts
- Provides hands-on, real world database experience using data from the City of Edmonton Open Data Portal
- Foster a collaborative workshop
 - Please interrupt and ask questions

Why SQL?

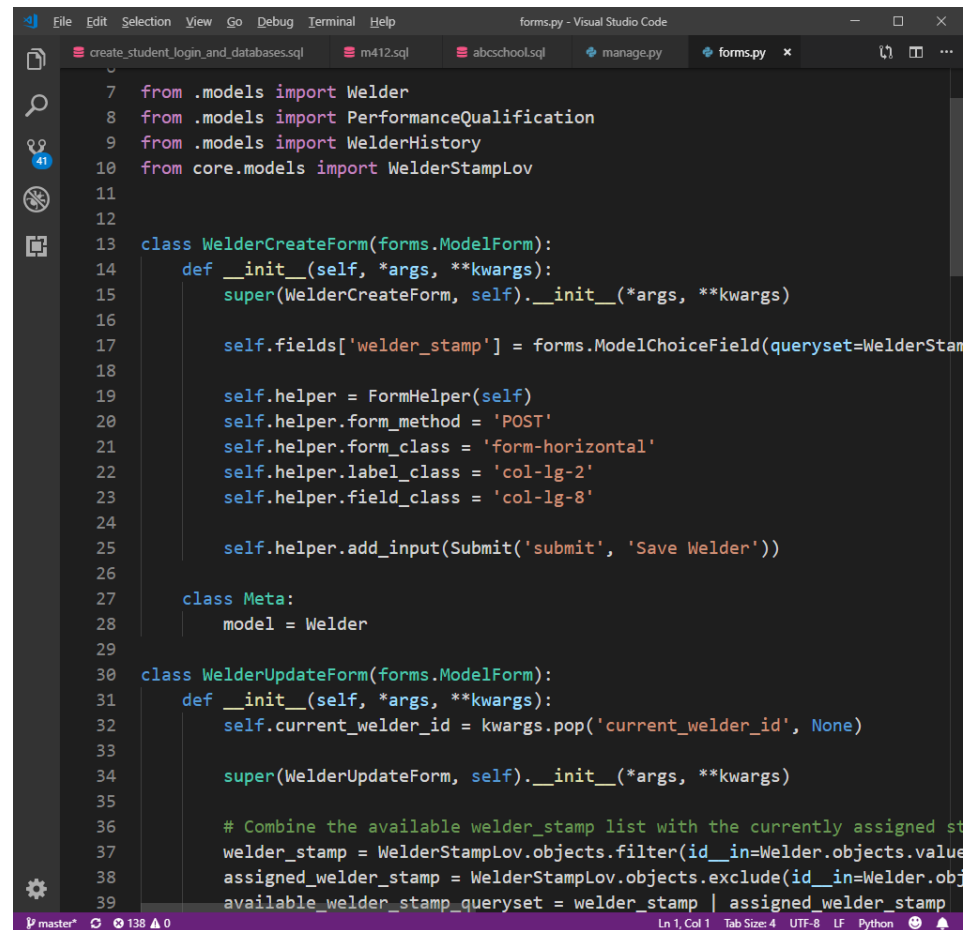
- Simple
- Accessible
- Applicable
- Powerful
- Pervasive
- Valuable
- Universal

```

88 );
89 CREATE TABLE SUBJECT_AREA (
90     SUBJECT_ID DECIMAL(2) NOT NULL,
91     SUBJECT_NAME VARCHAR(30) NOT NULL,
92     SUBJECT_EFF_DATE DATETIME NOT NULL,
93     SUBJECT_EXP_DATE DATETIME,
94     SUBJECT_TAX_PROFILE_CD VARCHAR(1) NOT NULL,
95     SUBJECT_CERTIFIED VARCHAR(1) NOT NULL,
96     SUBJECT_AUTH_LIMIT DECIMAL(9),
97     UPDATE_TS DATETIME NOT NULL,
98     PRIMARY KEY (SUBJECT_ID)
99 );
100 CREATE TABLE TRAINING_SITE (
101     SITE_ID DECIMAL(2) NOT NULL,
102     LOCATION VARCHAR(12) NOT NULL,
103     SITE_EFF_DATE DATETIME NOT NULL,
104     SITE_EXP_DATE DATETIME,
105     SITE_MAX_SIZE DECIMAL(3) NOT NULL,
106     UPDATE_TS DATETIME NOT NULL,
107     PRIMARY KEY (SITE_ID)
108 );
109
110 INSERT INTO ATTENDANCE(OFFERING_ID, STUDENT_ID, REGSTRN_MADE_DATE,
REGSTRN_CANC_DATE, EVALUATION, PENALTY, FINAL_MARK, AMOUNT_PAID, UPDATE_TS)
VALUES (9111, 98351, '2013-12-21 00:00:00.0', null, '7', '0', '7', 600,
'2014-01-05 18:14:15.0');
111 INSERT INTO ATTENDANCE(OFFERING_ID, STUDENT_ID, REGSTRN_MADE_DATE,
REGSTRN_CANC_DATE, EVALUATION, PENALTY, FINAL_MARK, AMOUNT_PAID, UPDATE_TS)
VALUES (9111, 84853, '2013-12-11 00:00:00.0', null, '6', '1', '5', 600,
'2015-01-05 19:14:15.0');
112 INSERT INTO ATTENDANCE(OFFERING_ID, STUDENT_ID, REGSTRN_MADE_DATE,
REGSTRN_CANC_DATE, EVALUATION, PENALTY, FINAL_MARK, AMOUNT_PAID, UPDATE_TS)
VALUES (9112, 98351, '2014-06-03 00:00:00.0', null, '5', '0', '5', 700,
'2014-07-01 10:14:15.0');
  
```


Why not Python? R?

- Difficult for beginners
- Complicated syntax
- Requires programming knowledge (logic, algorithms)
- Is SQL better than Python or R?
 - SQL is good for some things
 - Python/R is good for other things
 - Compliment each other
- SQL is a great starting point

A screenshot of a Visual Studio Code editor window titled 'forms.py - Visual Studio Code'. The editor shows Python code for Django forms. The code includes imports for 'Welder', 'PerformanceQualification', 'WelderHistory', and 'WelderStampLov' from '.models', and 'WelderStampLov' from 'core.models'. It defines two form classes: 'WelderCreateForm' and 'WelderUpdateForm', both inheriting from 'forms.ModelForm'. 'WelderCreateForm' has an '__init__' method that calls 'super().__init__()' and sets 'self.fields['welder_stamp']' to a 'ModelChoiceField' with 'queryset=WelderStampLov.objects'. It also sets 'self.helper' to 'FormHelper(self)' with 'form_method='POST'', 'form_class='form-horizontal'', 'label_class='col-lg-2'', and 'field_class='col-lg-8''. A 'Meta' class is defined with 'model = Welder'. 'WelderUpdateForm' has an '__init__' method that calls 'super().__init__()' and sets 'self.current_welder_id' to 'kwargs.pop('current_welder_id', None)'. It also has a comment and code to combine the available welder_stamp list with the currently assigned stamp. The status bar at the bottom shows 'master', '138', '0', 'Ln 1, Col 1', 'Tab Size: 4', 'UTF-8', 'LF', 'Python'.

Data Analytics

- Analytics is the discovery, interpretation, and communication of meaningful patterns in **data**; and the process of applying those patterns towards effective decision making
- Organizations may apply analytics to business **data** to describe, predict, and improve business performance
 - <https://en.wikipedia.org/wiki/Analytics>

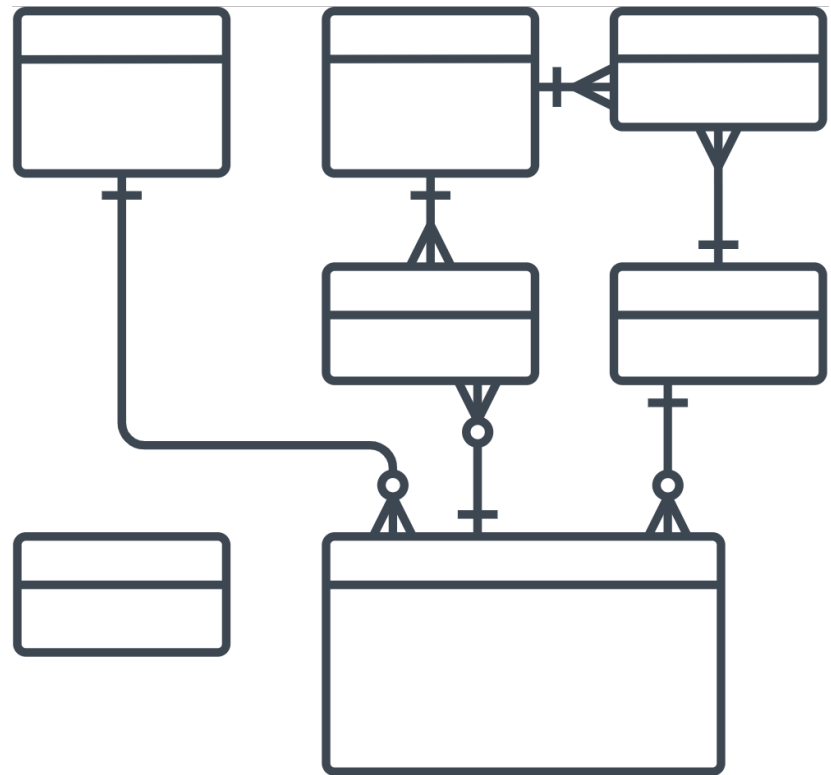
Relational Database

Workshop

Introducing SQL: Foundation of Data Analytics

What is a database?

- A relational “database” management system (RDBMS) organizes **data**
- The logical structure of the database is based upon the information needs of an organization
 - Entities (“things” of interest to the organization),
AND
 - Relationships (how the Entities are associated with each other)



Advantages of a RDBMS

- Establish a centralized, logical view of data
- Minimizes data duplication (i.e. “redundancy”)
- Promote data accuracy and integrity
- Capacity of database
- Superior multi-user or concurrent access
- Security
- Retrieve information quickly
- Inter-operability



<https://www.bespokesoftwaredevelopment.com/blog/advantages-database-development-business/>

Database Terminology

- **Table**, Entity, Relation, (similar to an Excel Worksheet)
- **Row**, Record, Instance
- **Column**, Field, Attribute
- **Primary Key** – unique and mandatory
- **Foreign Key** – a cross-reference between tables because it references the primary key of another table
- **Relationship** – created though foreign keys



How to introduce SQL?

- Microsoft Access
 - <https://products.office.com/en-ca/access>
- Microsoft SQL Server
 - <https://www.microsoft.com/en-us/sql-server/sql-server-2017>
- MariaDB, MySQL
 - <https://mariadb.org/>
 - <https://www.mysql.com/>
- Postgresql
 - <https://www.postgresql.org/>
- Oracle
 - <https://www.oracle.com/database/>
- Hadoop, Spark, Hive, Pig
 - <https://hadoop.apache.org/>



A database that ...

- Has full-featured SQL
- Has billions and billions of deployments
- Is a single-file database
- Has public domain source code
- Small footprint
- Has a max DB size of 140 terabytes
- Has a max row size of 1 gigabyte
- Is faster than direct file access
- Aviation-grade quality and testing
- Zero-configuration
- Has ACID (Atomic, Consistent, Isolated, and Durable) transactions, even after power loss
- Has a stable, enduring file format
- Is has extensive, detailed documentation
- Has long-term support (to the year 2050)

<https://www.sqlite.org/about.html>

SQLite

- “SQLite is the most widely deployed database in the world with more applications than we can count, including several high-profile projects”
 - <https://www.sqlite.org/famous.html>
- “SQLite is an in-process library that implements a self-contained, serverless, zero-configuration, transactional SQL database engine”
 - <https://www.sqlite.org/about.html>
- Perfect for learning SQL (the foundation of data analytics)

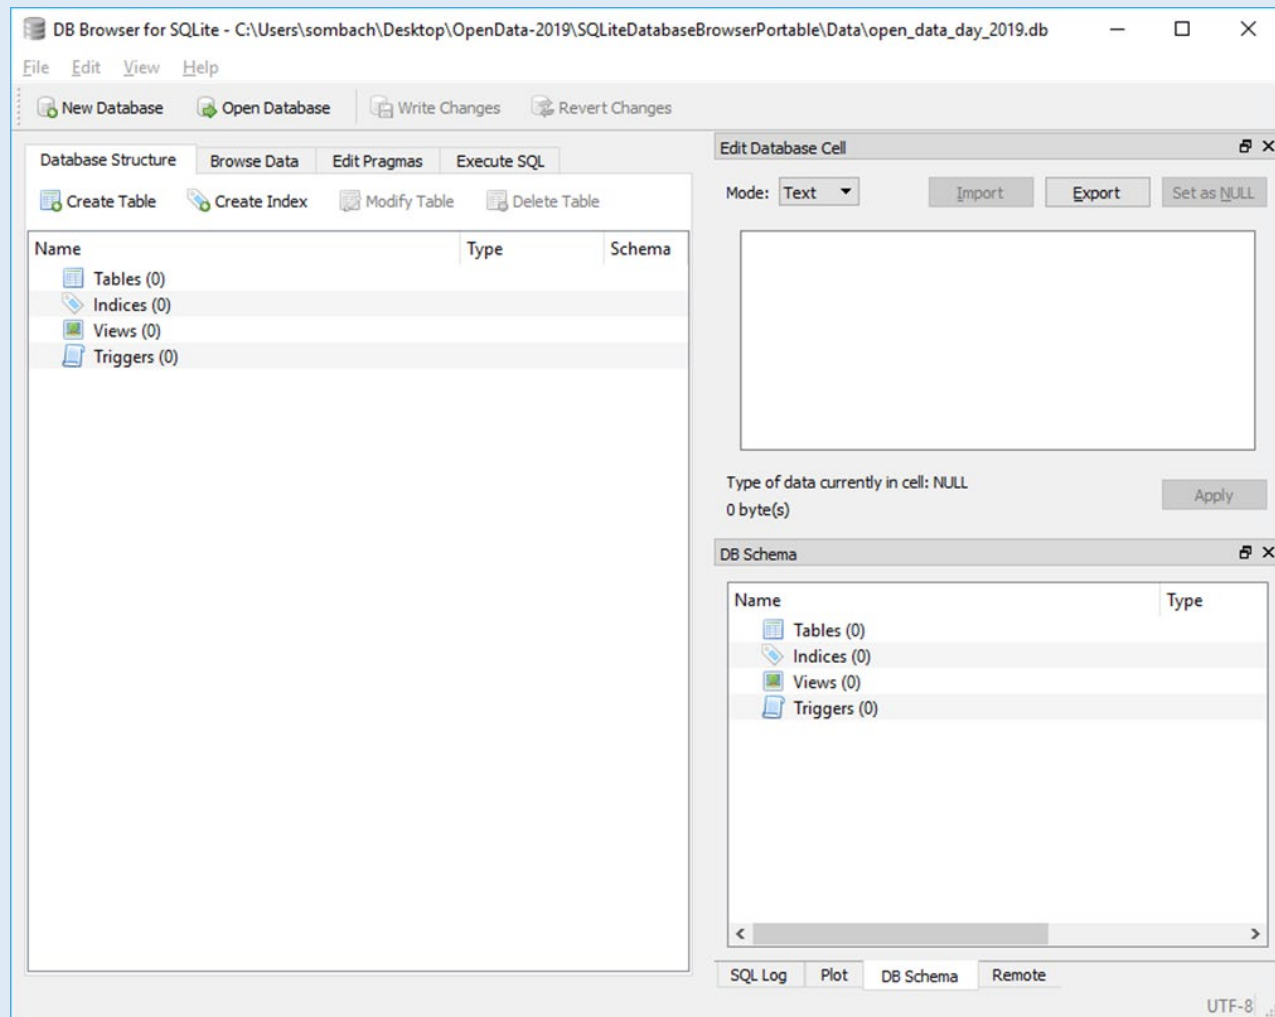
Exercise 1: Download and Run SQLite BD Browser

- Download SQLite
- Download SQLite DB Browser Portable
 - <https://sqlitebrowser.org/dl/>

Exercise 1: Download and Run SQLite

- Extract the ZIP archive to the Desktop
- Start SQLite
 - SQLiteDatabaseBrowserPortable.exe
- Create a New database
 - open_data_day_2019.db
- Save the database in the Data folder
- Click Cancel when prompted to create a table
- Done!

Exercise 1: Completed



SQL

Workshop

Introducing SQL: Foundation of Data Analytics

What is SQL?

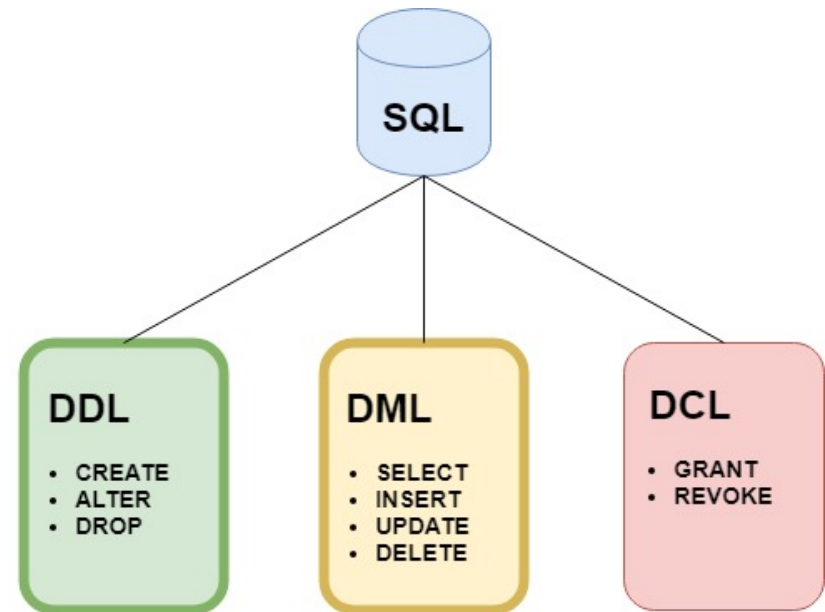
- SQL stands for Structured Query Language
 - SQL is pronounced S-Q-L or seque^l
 - SQL is a standard language for managing, manipulating and querying databases
 - Developed at IBM in the early 1970's
 - In 1986, ANSI and ISO standard groups officially adopted the standard "Database Language SQL" definition
 - Most SQL databases have their own proprietary extensions in addition to the SQL standard
- SQL is the language used to ask questions (query) of a database which will return answers (results)

Why is SQL the foundation of Data Analytics?

- Data engineers and database administrators will use SQL to ensure that everybody in their organization has access to the data they need
- Data scientists will use SQL to load data into their models
- Data analysts will use SQL to query tables of data and derive insights from it

Components of SQL

- SQL consists of three components which offer everything required to manage, maintain and use a database
 1. Data Definition Language
 2. Data Manipulation Language
 3. Data Control Language



Data Definition Language (DDL)

- This component is used to define the structure (or schema) of the database
- For tables there are three main commands:
 - **CREATE TABLE table_name**
 - To create a table in the database
 - **ALTER TABLE table_name**
 - To add or remove columns from a table in the database
 - **DROP TABLE table_name**
 - To remove a table from the database

Exercise 2: Data Definition Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the CREATE TABLE statement into the empty SQLite Execute SQL window
- Click the **Execute SQL** ► button on the toolbar
- If the table is created successfully, you should receive the following message:
 - Query executed successfully: CREATE TABLE "MOSQUITO_TRAP_DATA"
- Click Write Changes to make commit the changes permanent
- View the changes in the Database Structure tab

```
CREATE TABLE "MOSQUITO_TRAP_DATA" (  
  `SAMPLEID` INTEGER PRIMARY KEY AUTOINCREMENT,  
  `TRAP_DATE` NUMERIC,  
  `GENUS` TEXT,  
  `SPECIES` TEXT,  
  `TYPE` TEXT,  
  `GENDER` TEXT  
);
```

Exercise 2: Data Definition Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the ALTER TABLE statements into the empty SQLite Execute SQL window
- Click the **Execute SQL** ► button on the toolbar
- If the table is created successfully, you should receive the following message:
 - Query executed successfully: ALTER TABLE "MOSQUITO_TRAP_DATA"
- Click **Write Changes** to make commit the changes permanent
- View the changes in the **Database Structure** tab

```

ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RURALSOUTHWEST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RURALSOUTHEAST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RURALNORTHWEST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RURALNORTHEAST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RIVERVALLEYWEST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RIVERVALLEYEAST` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RESIDENTIALSOUTH` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `RESIDENTIALNORTH` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `LAGOON` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `GOLFCOURSE` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `INDUSTRIALPARK` INTEGER;
ALTER TABLE "MOSQUITO_TRAP_DATA" ADD COLUMN `TOTAL` INTEGER;

```

Exercise 2: Data Definition Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the DROP TABLE statement into the empty SQLite Execute SQL window
- Click the **Execute SQL** ► button on the toolbar
- If the table is created successfully, you should receive the following message:
 - Query executed successfully: DROP TABLE "MOSQUITO_TRAP_DATA"
- Click Write Changes to make commit the changes permanent
- View the changes in the Database Structure tab

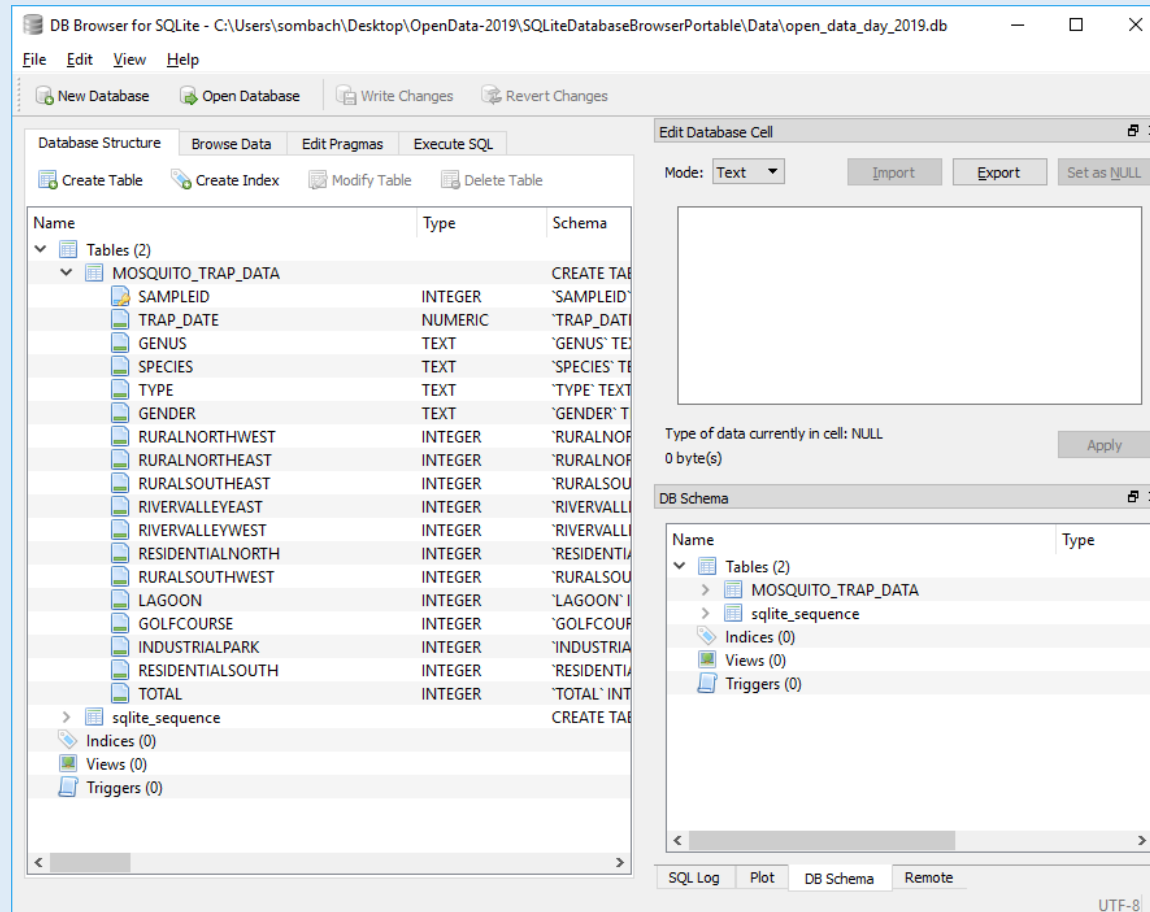
```
DROP TABLE "MOSQUITO_TRAP_DATA";
```

Exercise 2: Data Definition Language

- Create the MOSQUITO_TRAP_DATA table again using the DDL on the next slide
- Click **Write Changes** to make commit the changes permanent
- View the changes in the Database Structure tab
- Done!


```
CREATE TABLE "MOSQUITO_TRAP_DATA" (  
  `SAMPLEID` INTEGER PRIMARY KEY AUTOINCREMENT,  
  `TRAP_DATE` NUMERIC,  
  `GENUS` TEXT,  
  `SPECIES` TEXT,  
  `TYPE` TEXT,  
  `GENDER` TEXT,  
  `RURALSOUTHWEST` INTEGER,  
  `RURALSOUTHEAST` INTEGER,  
  `RURALSOUTHWEST` INTEGER,  
  `RIVERVALLEYEAST` INTEGER,  
  `RIVERVALLEYWEST` INTEGER,  
  `RESIDENTIALNORTH` INTEGER,  
  `RURALSOUTHWEST` INTEGER,  
  `LAGOON` INTEGER,  
  `GOLFCOURSE` INTEGER,  
  `INDUSTRIALPARK` INTEGER,  
  `RESIDENTIALSOUTH` INTEGER,  
  `TOTAL` INTEGER  
)
```

Exercise 1: Completed



Data Manipulation Language

- This component is used to manipulate data within a table
- There are four main commands:
 - SELECT
 - To select rows of data from a table
 - INSERT
 - To insert rows of data into a table
 - UPDATE
 - To change rows of data in a table
 - DELETE
 - To remove rows of data from a table

Exercise 3: SELECT

Data Manipulation Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the SELECT statement into the empty SQLite Execute SQL window
 - `SELECT COUNT(*) FROM MOSQUITO_TRAP_DATA;`
- Click the **Execute SQL** ► button on the toolbar
- Do you get an answer? Why not?

Exercise 3: INSERT

Data Manipulation Language

- Add some data to the MOSQUITO_TRAP_DATA table created in Exercise 2
- Type or copy/paste the INSERT statement into the empty SQLite Execute SQL window
- Click the **Execute SQL** ► button on the toolbar
- Click **Write Changes** to make commit the changes permanent
- View the changes in the **Browse Data** tab
- The MOSQUITO_TRAP_DATA table now has seven rows of data

```

INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','spencerii','Black legs','Female',0,0,0,0,0,1,0,0,0,1,1,3);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','dorsalis','Banded legs','Female',0,1,0,0,0,0,2,0,0,0,0,3);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','euedes','Banded legs','Female',1,1,0,0,2,0,0,0,0,0,0,4);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','excrucians','Banded legs','Female',1,2,0,0,2,1,0,0,0,1,0,7);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','fitchii','Banded legs','Female',0,2,0,0,1,0,0,0,0,0,0,4,7);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','flavescens','Banded legs','Female',6,5,8,0,0,0,5,0,0,3,1,28);
INSERT INTO "MOSQUITO_TRAP_DATA" (TRAP_DATE, GENUS, SPECIES, TYPE, GENDER, RURALNORTHWEST,
RURALNORTHEAST, RURALSOUTHEAST, RIVERVALLEYEAST, RIVERVALLEYWEST, RESIDENTIALNORTH,
RURALSOUTHWEST, LAGOON, GOLFCOURSE, INDUSTRIALPARK, RESIDENTIALSOUTH, TOTAL) VALUES ('2014-
07-01','Aedes','vexans','Banded legs','Female',3,168,1,21,38,8,16,0,0,3,32,290);

```

Exercise 3: SELECT

Data Manipulation Language

- Type or copy/paste the SELECT statement into the empty SQLite Execute SQL window
 - `SELECT COUNT(*) FROM MOSQUITO_TRAP_DATA;`
- Click the **Execute SQL** ► button on the toolbar
- When you execute the query, you are asking the database a question
 - Can you tell me the number of rows in the MOSQUITO_TRAP_DATA table?
- The database gives you an answer (the result) and you should have received the following message:
 - 7 rows returned in 1ms from: `SELECT * FROM MOSQUITO_TRAP_DATA;`

Exercise 3: SELECT

Data Manipulation Language

- What if you want to see all the rows in your database?
 - `SELECT * FROM MOSQUITO_TRAP_DATA;`
 - Returns all columns and rows in a table
- What if you only want to see the Genus, Species and Total of each row?
 - `SELECT GENUS, SPECIES, TOTAL FROM MOSQUITO_TRAP_DATA;`
 - Returns only the GENUS, SPECIES, TOTAL columns for each row in a table

Data Manipulation Language

- The WHERE clause
 - Uses operators to extract only those records that fulfill a specified condition
- Used to ask more complicated questions
- SQL will do exactly what you ask, not always what you expect
- “I do not think it means what you think it means”
 - *Inigo Montoya*

Operator	Description
=	Equal
<>	Not equal. Note: In some versions of SQL this operator may be written as !=
>	Greater than
<	Less than
>=	Greater than or equal
<=	Less than or equal
BETWEEN	Between a certain range
LIKE	Search for a pattern
IN	To specify multiple possible values for a column

Exercise 3: SELECT

Data Manipulation Language

- Show the rows that have a mosquito TYPE of “Black legs”
 - `SELECT * FROM MOSQUITO_TRAP_DATA WHERE TYPE = 'Black legs';`

YOUR TURN

- Write and execute a DML statement to answer the question below:
 - Which mosquito species' were caught in the traps placed in the west river valley?

Exercise 3: UPDATE

Data Manipulation Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the UPDATE statement into an empty SQLite Execute SQL window
- Click the **Execute SQL ▶** button on the toolbar
- You should receive the following message:
 - Query executed successfully: ... (took 1ms, 4 rows affected)

```
UPDATE MOSQUITO_TRAP_DATA  
SET GENDER = 'Male'  
WHERE SAMPLEID IN (1,3,5,7);
```

Data Manipulation Language

- The GROUP BY clause
 - Used in collaboration with the SELECT statement to arrange identical data into groups
- The GROUP BY statement is often used with aggregate functions

Function	Description
AVG	Calculates the average of a set of values
COUNT	Counts rows in a specified table or view
MAX	Gets the minimum value in a set of values
MIN	Gets the maximum value in a set of values
SUM	Calculates the sum of values

Exercise 3: SELECT

Data Manipulation Language

YOUR TURN

- Write and execute a DML statement to answer the question below:
 - How many mosquitos of each gender were caught in traps throughout the city?

Exercise 3: DELETE

Data Manipulation Language

- Select the Execute SQL tab in SQLite
- Type or copy/paste the DELETE statement into an empty SQLite Execute SQL window
- Click the **Execute SQL ►** button on the toolbar
- You should receive the following message:
 - Query executed successfully: ... (took 0ms, 4 rows affected)

```
DELETE FROM  
MOSQUITO_TRAP_DATA WHERE  
GENDER = "Male";
```


Exercise 3: SELECT

Data Manipulation Language

YOUR TURN

- Write and execute a DML statement to answer the question below:
 - At which traps were more mosquitos caught? Rural north east or rural north west?
- Done!

Advanced SQL

- The MOSQUITO database only has one table
- Databases with more than one table require tables to be joined
- Foreign keys create relationships between tables and must be joined in a DML statement

- Download the LED Streetlight Conversion database called odd_streetlight.db
- Execute the query below

```
SELECT LED_STREETLIGHT.STREETLIGHT_ID, LED_STREETLIGHT.TYPE,  
LOCATION.LOCATION  
FROM LED_STREETLIGHT, LOCATION  
WHERE LED_STREETLIGHT.STREETLIGHT_ID = LOCATION.STREETLIGHT_ID  
AND LED_STREETLIGHT.STREETLIGHT_ID = 12;
```

City of Edmonton Open Data Portal

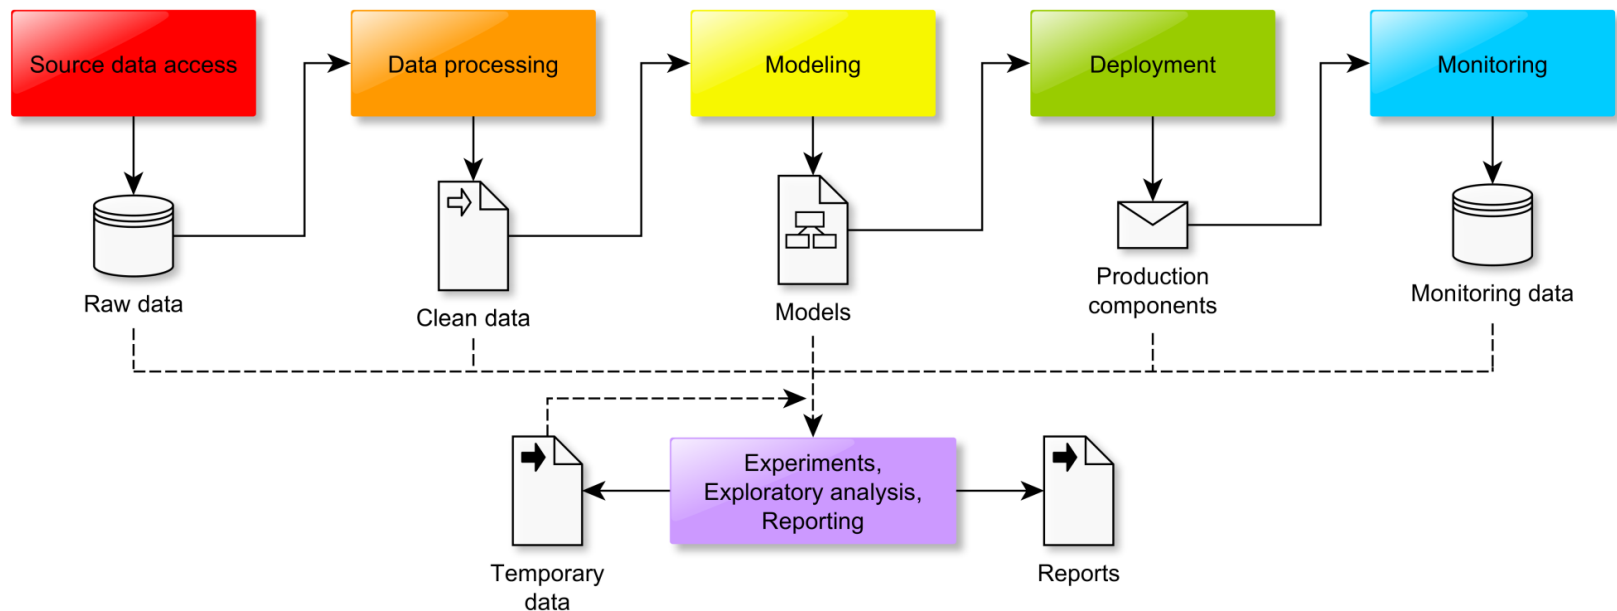
Workshop

Introducing SQL: Foundation of Data Analytics

Using the Open Data Portal

- <https://data.edmonton.ca/>
- Data sets are usually available in comma separated value (CSV) format
- To use the dataset requires cleaning, importing, exploring and understand the data set
 - Workshop: Exploring & Cleaning Data with OpenRefine
- Requires work

Data Work Flow



How I prepared the data sets for today

- Selected data sets from the Open Data Portal
- Downloaded the CSV and surveyed in Google Sheets
- Cleaned the data set
 - E.g. reformatted dates from MMM DD YYYY to YYYY-MM-DD
- Imported into directly into SQLite tables
- Added primary keys
- Explored data set using DML

Some “Mosquitoes Trap Data” questions

- How many mosquitos caught in 2014?
SELECT strftime('%Y', TRAP_DATE) as YEAR, SUM(TOTAL)
FROM MOSQUITO_TRAP_DATA
WHERE TOTAL <> ''
AND TOTAL > 0
GROUP BY YEAR;
- How many mosquitos of each species were caught?
- Which traps caught the most mosquitos?

Some “LED Streetlight Conversion” questions

- How many total streetlights?
- How many streetlights are converted to LED?
- How many streetlights were converted by year?

```
SELECT strftime('%Y', STARTDATE) as YEAR, TYPE,  
COUNT(STREETLIGHT_ID)  
FROM LED_STREETLIGHT  
WHERE TYPE = "LED"  
GROUP BY YEAR;
```

SQL and Climate Change

- Connecting and linking various data sets
- Builds an understanding of what that data means
- Data is a universal language,
climate change is a global
problem

Next steps

- Playing with data and SQL forces you to think and understand the data (builds knowledge)
 - The relationships between data
 - The meaning of those relationships
 - The validity of the data
- SQL is iterative, often a “trial and error” process
 - Don’t be afraid to make mistakes
 - Team sport – discuss, share, question, collaborate
- Data is everywhere which raises questions of privacy, security and ethics

Experiment



<https://www.manchester.ac.uk/discover/news/major-leap-towards-storing-data-at-the-molecular-level/>

If there's time ... (I talked too fast)

- Let's (democratically):
 1. Choose a dataset not discussed during the workshops
 2. Formulate a question related to the dataset
 3. Load the data into SQLite
 4. Execute some DML to answer the question

Thank you!

- Robb Sombach
 - sombach@ualberta.ca
 - robb@sombach.com
 - [LinkedIn](#)



References

- <https://opendataday.org/>
- <https://data36.com/sql-for-data-analysis-tutorial-beginners/>
- <https://www.datascience.com/blog/to-sql-or-not-to-sql-that-is-the-question>
- <https://codebeautify.org/sqlformatter>